

**IMPORTANT** Please remember to destroy all the resources after each work session. You can recreate infrastructure by creating new PR and merging it to master.

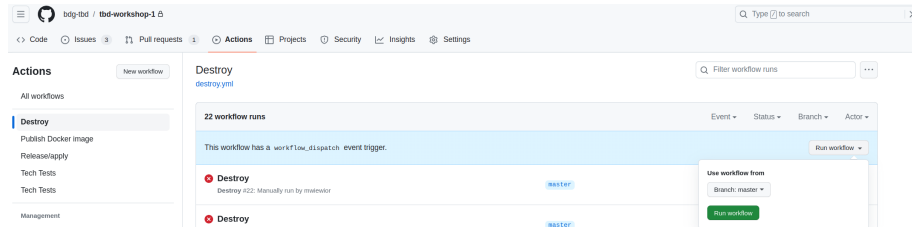


Figure 1: img.png

0. The goal of this phase is to create infrastructure, perform benchmarking/scalability tests of sample three-tier lakehouse solution and analyze the results using:

- TPC-DI benchmark
- dbt - data transformation tool
- GCP Composer - managed Apache Airflow
- GCP Dataproc - managed Apache Spark
- GCP Vertex AI Workbench - managed JupyterLab

Worth to read: \* <https://docs.getdbt.com/docs/introduction> \* <https://airflow.apache.org/docs/apache-airflow/stable/index.html> \* <https://spark.apache.org/docs/latest/api/python/index.html> \* <https://medium.com/snowflake/loading-the-tpc-di-benchmark-dataset-into-snowflake-96011e2c26cf> \* <https://www.databricks.com/blog/2023/04/14/how-we-performed-etl-one-billion-records-under-1-delta-live-tables.html>

2. Authors:

- Zuzanna Górecka
- Adam Górski
- Michał Oracki

Zespół nr 7

***Link to forked repo***

3. Sync your repo with <https://github.com/bdg-tbd/tbd-workshop-1.0>.

4. Provision your infrastructure.

- a) setup Vertex AI Workbench **pyspark** kernel as described in point 8
- b) upload tpc-di-setup.ipynb to the running instance of your Vertex AI Workbench

5. In **tpc-di-setup.ipynb** modify cell under section ***Clone tbd-tpc-di repo:***

a)first, fork <https://github.com/mwiewior/tbd-tpc-di.git> to your github organization.

b)create new branch (e.g. 'notebook') in your fork of tbd-tpc-di and modify profiles.yaml by commenting following lines:

```
#"spark.driver.port": "30000"
#"spark.blockManager.port": "30001"
#"spark.driver.host": "10.11.0.5" #FIXME: Result of the command (kubectl get node
#"spark.driver.bindAddress": "0.0.0.0"
```

This lines are required to run dbt on airflow but have to be commented while running dbt in notebook.

c)update git clone command to point to *your fork*.

6. Access Vertex AI Workbench and run cell by cell notebook `tpc-di-setup.ipynb`.

a) in the first cell of the notebook replace: `%env DATA_BUCKET=tbd-2023z-9910-data` with your data bucket.

b) in the cell: 

```
%%bash      mkdir -p git && cd git      git
clone https://github.com/mwiewior/tbd-tpc-di.git      cd
tbd-tpc-di      git pull replace repo with your fork. Next
checkout to 'notebook' branch.
```

c) after running first cells your fork of `tbd-tpc-di` repository will be cloned into Vertex AI enviroment (see git folder).

d) take a look on `git/tbd-tpc-di/profiles.yaml`. This file includes Spark parameters that can be changed if you need to increase the number of executors and

```
server_side_parameters:
  "spark.driver.memory": "2g"
  "spark.executor.memory": "4g"
  "spark.executor.instances": "2"
  "spark.hadoop.hive.metastore.warehouse.dir": "hdfs:///user/hive/warehouse/"
```

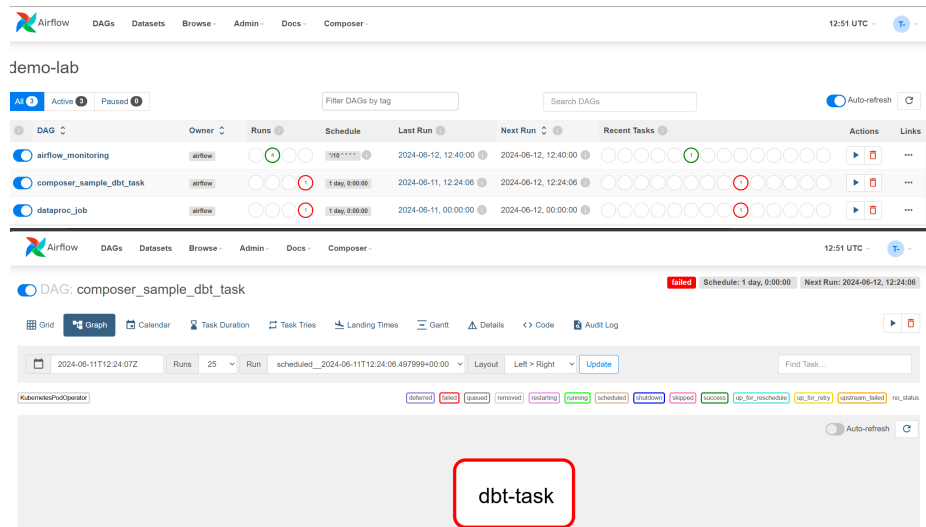
7. Explore files created by generator and describe them, including format, content, total size.

There are 3 directories Batch1 Batch2 and Batch3 with each of them containing multiple csv and text files. The biggest ones are DailyMarket.txt and WatchHistory.txt with sizes of 296MB and 134MB. Directioes contain also multiple FINWIRE files varying in size from 900KB to 70KB.

8. Analyze `tpcdi.py`. What happened in the loading stage?

During loading as no `file_name` is specified - all are read and saved to the bucket which has been set up as an environment variable. This includes: `DATE`, `DAILY_MARKET`, `INDUSTRY`, `PROSPECT`, `CUSTOMER_MGMT`, `TAX_RATE`, `HR`, `WATCH_HISTORY`, `TRADE`,





Analyzing the logs (including in Airflow), we concluded that there is a problem with memory limits. Despite many attempts, we were unable to resolve it. We tried changing the machine type and memory limits.

However, we were the first team to solve the problem that occurred during release at the beginning of task 2a. Thanks to this, other teams could solve the problem using our recommendations. After analyzing the logs and tracking CPU and SSD usage, we noticed that the problem was with goutas. The problem was that the cluster tried to autoscale but couldn't due to upper limits, so after a few failed attempts it deleted itself and healthchecks had no way to reach the required pods.